

DATAACENDIA

AI Red-Teaming Design Partnership

- Adversarial testing of AI decisions — CendiaApotheosis and CendiaRedTeam
- Decision Crisis Immunization Infrastructure

"Your AI gets stronger while you sleep. CendiaApotheosis runs nightly red-teaming, auto-patching, and continuous upskilling."

— The Datacendia Manifesto

WHAT WE BUILD

Decision Crisis Immunization Infrastructure

- Multi-agent deliberation — 50+ specialized agent presets
- Cryptographic audit trail — immutable evidence
- Zero-copy architecture — data never leaves your DB
- Sovereign deployment — air-gapped to cloud
- Not a chatbot. A system of specialized intelligences that debate, dissent, and decide — with cryptographic proof of every step.
- 50+ agent presets (CFO, CISO, Risk, Strategy, Legal) deliberate on your data. You get a boardroom debate, not a chatbot answer.
- Every decision, vote, dissent — hashed, timestamped, exportable. When the auditor says "prove it," hand them this.

AI RED-TEAMING

Red-Teaming Capabilities

- Adversarial challenge of every recommendation
- Nightly red-teaming with auto-patching
- Formal dissent with retaliation protection

USE CASES

Red-Teaming Use Cases

- Decision stress-testing — Adversarial challenge of critical recommendations
- AI bias detection — Systematic testing for bias in AI outputs
- Security assessment — Red-teaming of AI security posture
- Compliance testing — Adversarial testing against regulatory requirements
- Model validation — Independent challenge of AI model recommendations

DESIGN PARTNER BENEFITS

Why Partner Early?

- Direct influence on product roadmap
- Priority support and dedicated engineering
- Preferred pricing locked at design partner rates
- Co-marketing — joint case studies
- Early access to new features
- 90-day pilot at reduced rate
- Weekly syncs with engineering
- Feature input on quarterly roadmap

PLATFORM METRICS

156

Validated Endpoints

16

Connectors

26

Languages


4

Deployment Modes

MULTI-AGENT DECISION ENGINE

Decisions Made by Specialized AI Agents

- Role-specific agents: CFO, CISO, Legal, Strategy, Risk
- Confidence scores and reasoning — dissent preserved
- War Room, Governed, and Quick Brief modes



Every Decision Tracked with Full Evidence Chain

- Complete searchable decision history
- Full deliberation records linked to each decision
- Trace any decision back to evidence in seconds

←

All Decisions

6 decisions found

Ask Council

Refresh

Q Search decisions...

▽ All Modes

▽ All Status

Should we expand into the European market in Q2?

Deliberation

Completed

0.87% confidence

6 agents

Mar 15, 2026, 08:14 PM

The Council recommends proceeding with EU expansion, prioritizing Germany and Netherlands. Risk-adjusted ROI projects 34% within 18 months. Key risks: regulatory compliance (GDPR Article 22), currency exposure, and local hiring timeline.

Evaluate the acquisition of DataSync Corp for \$12M

Deliberation

Completed

0.72% confidence

8 agents

Mar 14, 2026, 08:14 PM

Split recommendation: CFO and COO agents favor acquisition for talent and IP. CISO agent flagged 3 critical security vulnerabilities in target codebase. Legal agent noted pending IP litigation. Recommend extending due diligence by 30 days.

What is the optimal pricing strategy for Enterprise tier?

Deliberation

Completed

0.91% confidence

5 agents

Mar 13, 2026, 08:14 PM

Strong consensus on value-based pricing at \$2,400/seat/year with volume discounts (10+ seats: 15%, 50+: 25%). Projects 28% increase in enterprise pipeline conversion. CMO agent recommends bundling compliance module as differentiator.

Risk assessment for migrating to multi-region cloud infrastructure

Quick Brief

Completed

0.83% confidence

7 agents

Mar 12, 2026, 08:14 PM

Red team identified 3 critical risks: data sovereignty gaps in APAC region, 340ms latency spikes during cross-region failover, and \$180K/month cost increase. Mitigation plan provided with phased rollout starting US-East → EU-West → APAC.