

SERIES A

Decision CrisisImmunizationInfrastructure

→ 50+ specialized AI agent presets debate, dissent, and decide — with cryptographic proof of every step. Not a chatbot. A boardroom.

THE PROBLEM

AI Makes Decisions. Nobody Can Explain Them.

- \$7.9T in decisions made using AI by 2030 with no audit trail
- 73% of enterprises cite AI explainability as top concern (Gartner)
- \$2.1B in regulatory fines for unexplainable algorithmic decisions (2023)
- Zero platforms provide multi-agent deliberation with cryptographic proof
- Enterprises are deploying AI at scale — but every decision is a black box. When the regulator asks "why did your AI do that?", the answer is silence.
- Single-model AI gives you one answer. Real decisions require debate, dissent, and documented reasoning. Boards don't accept one person's opinion — why accept one model's?
- Datacendia fills the gap between "AI generated an answer" and "AI made a defensible decision."

The Council™

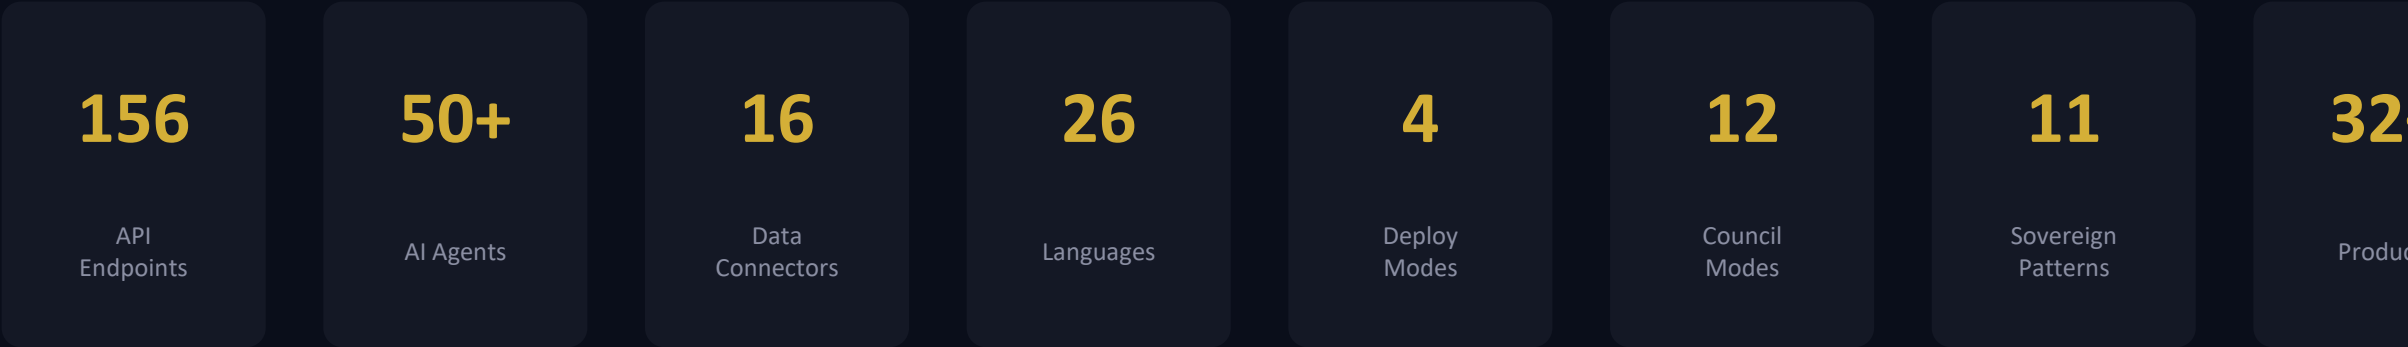
- 50+ specialized AI agent presets with distinct expertise deliberate on your data — like a boardroom, not a chatbot.
- CendiaChronos (time machine), Ghost Board, CendiaPreMortem, Decision Debt, CendiaLive, CendiaReplay, CendiaEcho, CendiaLens, CendiaCollapse, CendiaPulse, CendiaCrisis, CendiaROI, CendiaDCII.
- CendiaOversight, CendiaNotary, CendiaVault, CendiaProvenance, Decision Packets, CendiaCrucible, CendiaRedTeam, SGAS, CendiaCourt, CendiaSandbox, CendiaZKP, CendiaInsure, CendiaQuantumKMS, CendiaCarbon, CendiaJurisdiction, CendiaCompliance.
- Cloud to air-gapped. 4 deployment modes. 11 sovereign patterns. Zero-copy architecture — your data never leaves your database.
- 16 native connectors. Data stays in your DB — Datacendia queries in place. No ETL, no data movement, no risk.
- CendiaOmniTranslate (100 languages), CendiaDissent (protected whistleblower channel), CendiaResponsibility (human accountability layer with TPM-signed liability).
- Every tool purpose-built for AI decision governance. From real-time deliberation (CendiaLive) to quantum-resistant signing (CendiaQuantumKMS) — no competitor offers this depth.

THE COUNCIL™ IN ACTION

Multi-Agent Deliberation — Live UI

PLATFORM SCALE

Built for Enterprise, Engineered for Rigor



THE COUNCIL™

50+ Specialized Agent Presets

- Security & Compliance
- Adversarial Challenge
- Each agent brings domain expertise, challenges assumptions, and votes with confidence scores. Dissenting opinions are preserved — not silenced. The Arbiter synthesizes all perspectives into an actionable decision with full evidence.

DECISION OUTPUT

Auditable Decision Package

Deploy Anywhere — Your Data, Your Control

- Data Diode, Local RLHF, Shadow Council, TPM Attestation, Federated Mesh, QR Air-Gap Bridge, and more.
- Data never moves. Datacendia queries in place via native connectors. No ETL pipelines, no data lakes to maintain.
- Architecture designed to FedRAMP Moderate controls. SOC 2 Type II In Progress. ISO 27001 target H2 2026.

PRICING

Transparent, Procurement-Ready

- Implementation (\$15K–\$120K), Consulting (\$8K–\$25K), Training (\$2.5K–\$8K), Custom Development (\$15K–\$100K), Premium Support (\$2.5K–\$6K/mo).
Average services attach rate: 30–40% of ACV.
- Gross margin >80%. Net dollar retention >130%. CAC payback <18 months. LTV:CAC >5x. Land with Pilot, expand to Enterprise.

Category-Creating in a \$50B+ Market

- EU AI Act — mandates explainability for high-risk AI systems (2026 enforcement)
- US Executive Order on AI — federal agencies must document AI decision processes
- DORA (Digital Operational Resilience Act) — financial services AI governance mandate
- Enterprise AI spend — projected \$200B+ by 2028, governance is lagging
- Regulation is arriving — first-mover advantage in compliance tooling
- LLMs made multi-agent possible — viable at enterprise scale since 2024
- Enterprise AI budgets unlocked — every F500 has an AI governance gap
- Trust deficit growing — AI hallucinations eroding confidence

COMPETITIVE MOAT

Seven Layers of Defensibility

- 50+ specialized agent presets with domain reasoning — not wrapper-on-ChatGPT
- Immutable hash chain for every decision — competitors can't retrofit this
- Air-gapped to cloud — 11 sovereign patterns no competitor offers
- Data never moves — eliminates entire attack surface and compliance burden
- Day-one global enterprise readiness with RTL support
- Full audit theater — replay any decision step by step for any stakeholder
- More decisions → better agents → more accurate deliberations → higher switching cost
- Bottom line: Traditional BI tools (Tableau, Power BI) show you data. LLM wrappers (ChatGPT Enterprise) give one opinion. Palantir moves your data. Only Dacendia provides multi-agent deliberation with cryptographic proof, zero-copy architecture, and sovereign deployment — simultaneously.

PLATFORM PROOF

Production-Ready — Not a Prototype

COMPLIANCE

Enterprise Compliance Roadmap

- Why this matters for fundraising: Every compliance framework unlocks a market segment. SOC 2 → enterprise SaaS. HIPAA → healthcare (\$4T market). FedRAMP → federal government (\$100B+ IT spend). Each certification is a sales accelerant, not just a checkbox.

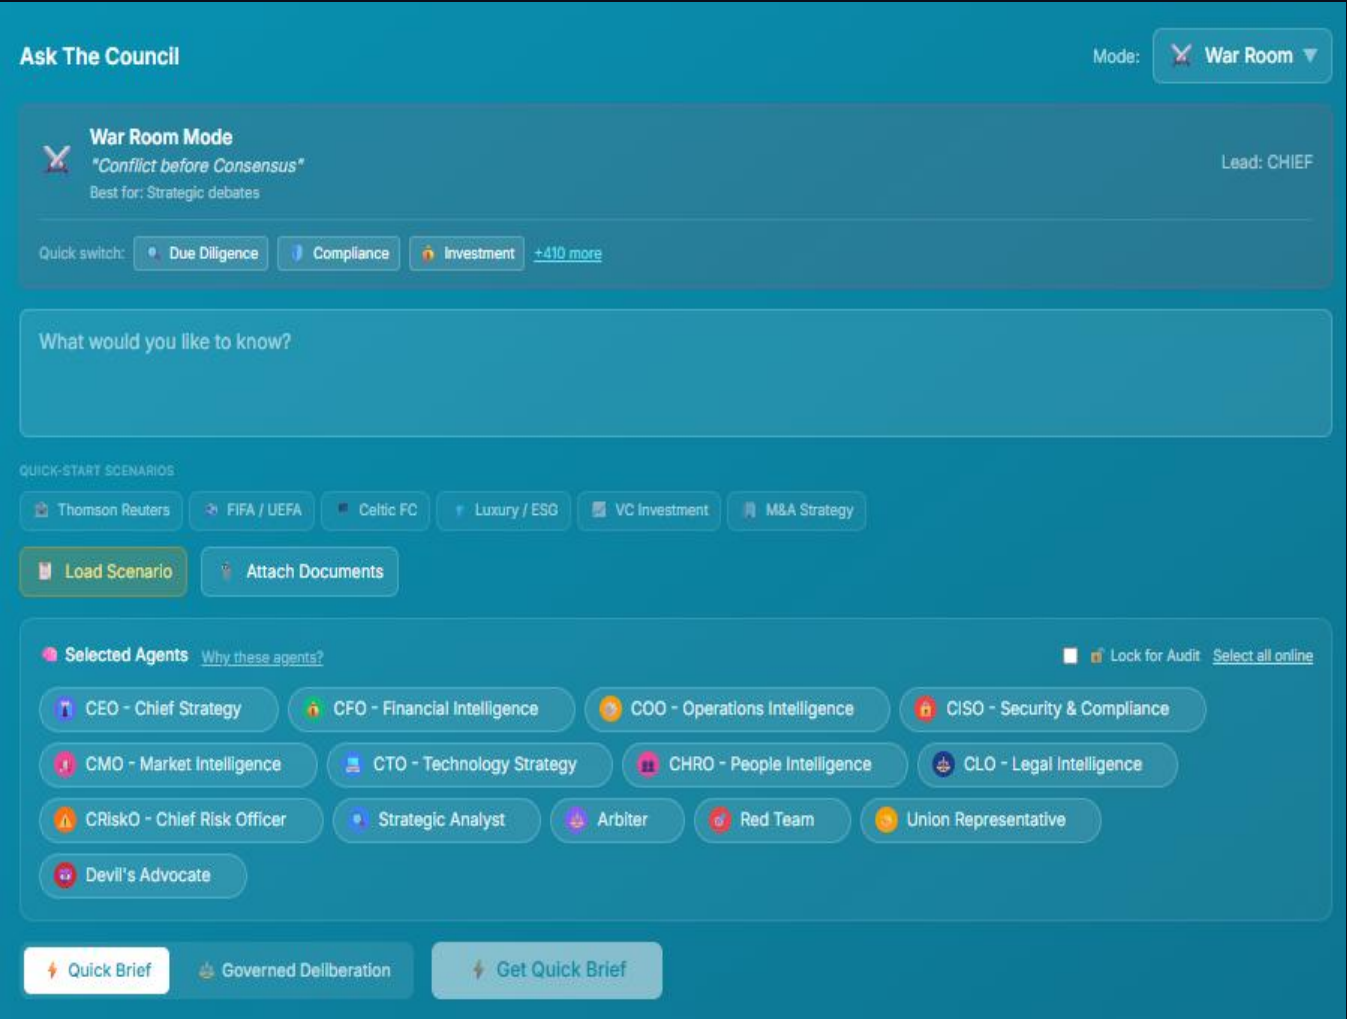
"Deploy Datacendia for 90 days on a single strategic problem. If the Council doesn't identify at least 3 critical risks you missed, we refund the pilot fee."

— The Datacendia Manifesto

MULTI-AGENT DECISION ENGINE

Decisions Made by Specialized AI Agents

- Role-specific agents: CFO, CISO, Legal, Strategy, Risk
- Confidence scores and reasoning — dissent preserved
- War Room, Governed, and Quick Brief modes



Every Decision Tracked with Full Evidence Chain

- Complete searchable decision history
- Full deliberation records linked to each decision
- Trace any decision back to evidence in seconds

© 2024-2026 Datacendia, LLC