

DEFENSE & GOVERNMENT

Sovereign AI Decision Infrastructure

→ Air-gapped multi-agent deliberation with cryptographic audit trails. Built for classified environments, designed for mission-critical decisions.

THE CHALLENGE

AI Decisions in Classified Environments Have Zero Accountability Infrastructure

- No audit trail — AI recommendations in C4ISR have no cryptographic proof chain
- Single-model risk — one model's hallucination becomes mission-critical error
- Data sovereignty mandated — cloud AI violates classification requirements
- Explainability required — DoD Responsible AI Strategy demands it
- Speed vs rigor trade-off — current tools force one or the other
- Responsible AI — explainable, governable, auditable
- Human-AI teaming — AI assists, humans decide
- Data-centric AI — decisions grounded in evidence

11 Sovereign Patterns — Air-Gapped to Cloud

- One-way data flow. Analysis results out, no data in.
- On-prem model fine-tuning. No external training data leakage.
- Parallel classified + unclassified deliberation for comparison.
- Hardware-level integrity verification for every compute node.
- Multi-site deployment with no central data aggregation.
- Optical data transfer between air-gapped networks.
- SGX/SEV hardware enclaves for computation.
- Prove decision validity without revealing inputs.

PLATFORM IN ACTION

Multi-Agent Deliberation — Live UI

MISSION APPLICATIONS

Decision Intelligence Across the Mission Stack

- Key differentiator: Every deliberation produces a cryptographic hash chain — immutable evidence of who (which agent) said what, when, with what confidence, and what dissents were raised. This is the audit infrastructure the DoD Responsible AI Strategy demands.

Built for Government Compliance

- SBIR/STTR — Phase I/II via DARPA, NSF, AFWERX
- OTA (Other Transaction Authority) — via DIU prototype contracts
- GSA Schedule — target listing H2 2026
- Direct Contract — capability statements available
- Partner Channel — via Booz Allen, GDIT, Leidos, SAIC

ARCHITECTURE

Zero-Trust, Zero-Copy, Zero-Egress

156

API Endpoints

50+

AI Agents

16

Connectors

11

Sovereign Patterns

"Deploy Datacendia on a single mission-relevant problem for 90 days. If the Council doesn't identify at least 3 critical risks that your current analysis missed, we refund the pilot fee."

— The Datacendia Manifesto

MULTI-AGENT DECISION ENGINE

Decisions Made by Specialized AI Agents

- Role-specific agents: CFO, CISO, Legal, Strategy, Risk
- Confidence scores and reasoning — dissent preserved
- War Room, Governed, and Quick Brief modes



